

# Practical Work 4: Advanced Vision, Segmentation, and 3D Data

Louis Fippo Fitime, Claude Tinku, Kerolle Sonfack  
Department of Computer Engineering, ENSPY

October 8, 2025

## Abstract

This fourth Practical Work focuses on advanced computer vision tasks that demand high precision and architectural complexity. Students will apply Deep Learning Engineering best practices to implement and evaluate a cutting-edge **U-Net architecture** for **semantic segmentation of medical images**. Furthermore, the module introduces the theoretical and practical challenges of handling **3D volumetric data** using Conv3D layers.

## Learning Objectives

- Master the **semantic segmentation** task and the use of the **U-Net** architecture.
- Understand and implement specific metrics for segmentation (**IoU**, **Dice Coefficient**).
- Systematically apply **Deep Learning Engineering** best practices (experiment tracking, model packaging).
- Understand the core concepts and implementation of **3D Convolutions** for volumetric data.
- Recognize the specific challenges of medical image data (imbalance, limited size).

# 1 Part 1: Segmentation and MLOps Best Practices (1.5h)

## 1.1 Semantic Segmentation and the U-Net Architecture

Semantic segmentation is the task of assigning a class label to every pixel in an image, which is crucial in fields like autonomous driving and medicine.

- **Output Type:** Unlike classification, what is the dimension and nature of the output tensor of a semantic segmentation model?
- **U-Net Structure:** The U-Net is famous for its "U" shape and its use of **skip connections** between the contracting (encoder) and expansive (decoder) paths. Explain the role of the decoder path and how skip connections differ here compared to standard ResNet blocks.
- **Loss Functions:** Why is the standard categorical cross-entropy often inadequate for medical segmentation tasks where the foreground (e.g., tumor) is tiny compared to the background? Propose an alternative loss function.

## 1.2 Engineering Practices: Experiment Tracking

In advanced DL, reliable comparison between experiments is mandatory. We will continue to use MLflow (introduced in TP 2).

### Instructions:

1. Define a strict naming convention for your MLflow runs to clearly identify the architecture, optimizer, and loss function used for each segmentation model trained.
2. Review the process of logging **custom metrics** (like the Dice Coefficient) which are not native to Keras's basic metrics list.

## 2 Part 2: Semantic Segmentation on Medical Data (3h)

### 2.1 Exercise 1: Implementing the U-Net Architecture

You will build a simplified 2D U-Net for segmenting a common medical dataset (e.g., a simulated cell or organ mask task). Assume the data has been pre-processed into normalized image arrays X and corresponding binary mask arrays Y.

**Instructions:**

1. Create a function for the core **convolutional block** (Conv-Batch Norm-ReLU).
2. Implement the **U-Net Encoder Path** using Max Pooling for downsampling.
3. Implement the **U-Net Decoder Path** using Conv2D Transpose (Upsampling).
4. **Crucial Step:** Implement the skip connections (concatenation) between the encoder and decoder.

```

1 import tensorflow as tf
2 from tensorflow import keras
3
4 def conv_block(input_tensor, num_filters):
5     # Core convolutional block
6     x = keras.layers.Conv2D(num_filters, (3, 3), padding='same')(input_tensor)
7     x = keras.layers.BatchNormalization()(x)
8     x = keras.layers.Activation('relu')(x)
9
10    x = keras.layers.Conv2D(num_filters, (3, 3), padding='same')(x)
11    x = keras.layers.BatchNormalization()(x)
12    x = keras.layers.Activation('relu')(x)
13
14    return x
15
16 def build_unet(input_shape=(128, 128, 1)):
17     inputs = keras.Input(input_shape)
18
19     # ENCODER PATH (Contracting)
20     c1 = conv_block(inputs, 32)
21     p1 = keras.layers.MaxPooling2D((2, 2))(c1)
22
23     # TODO: Implement 2 more contracting steps (c2, p2 and c3, p3)
24     # The number of filters should double at each step (64, 128)
25
26     # BRIDGE / BOTTLENECK (256 filters)
27     b = conv_block(p3, 256)
28
29     # DECODER PATH (Expansive)
30     # Step 1: Upsampling + Skip Connection
31     u1 = keras.layers.Conv2DTranspose(128, (2, 2), strides=(2, 2),
32                                     padding='same')(b)
33
34     # TODO: Concatenate the corresponding encoder output (c3) with u1

```

```

32 u1 = keras.layers.concatenate([u1, c3])
33 d1 = conv_block(u1, 128)
34
35 # TODO: Implement 2 more expansive steps (u2, d2 and u3, d3)
36 # ...
37
38 # Output Layer: 1 filter (for binary segmentation) with sigmoid
# activation
39 outputs = keras.layers.Conv2D(1, (1, 1), activation='sigmoid')(d3)
40
41 return keras.Model(inputs=[inputs], outputs=[outputs])
42
43 TODO: Compile and train the model using a suitable loss and custom
metric (IoU/Dice)
44 model.compile(...)

```

Listing 1: Simplified 2D U-Net Implementation with Keras

## 2.2 Exercise 2: Segmentation-Specific Metrics

The Dice Similarity Coefficient (or F1-Score for segmentation) is a common metric. It is defined as:

$$\text{Dice} = \frac{2 \cdot |A \cap B|}{|A| + |B|}$$

where A is the predicted mask and B is the ground truth mask.

**Instructions:**

1. Implement the Dice Coefficient as a custom Keras metric function.
2. Implement the IoU (Intersection over Union) metric.
3. **Question:** Compare the sensitivity of the IoU metric versus the Dice Coefficient to small segmentation errors, especially when the target mask is small.

```

1 from keras import backend as K
2
3 def dice_coeff(y_true, y_pred, smooth=1.):
4     # Flatten the tensors for calculation
5     y_true_f = K.flatten(y_true)
6     y_pred_f = K.flatten(y_pred)
7
8     # Calculate intersection and union
9     intersection = K.sum(y_true_f * y_pred_f)
10
11    # TODO: Complete the Dice formula calculation
12    # return ...
13
14    # Example for IoU
15    def iou_metric(y_true, y_pred):
16        intersection = K.sum(K.abs(y_true * y_pred), axis=-1)

```

```
17     union = K.sum(y_true, axis=-1) + K.sum(y_pred, axis=-1) -  
18         intersection  
    return K.mean((intersection + smooth) / (union + smooth), axis  
                  =-1)
```

Listing 2: Dice Coefficient Metric Implementation

### 3 Part 3: Introduction to 3D Convolutions and Volumetric Data (2.5h)

#### 3.1 Concepts: Conv3D for Volumetric Data

Medical imaging often involves volumetric data (e.g., CT scans, MRI) which are stacks of 2D slices, forming a  $D \times H \times W \times C$  tensor (Depth, Height, Width, Channels). Standard Conv2D layers only process  $H \times W$  spatial dimensions, ignoring the  $D$  dimension.

- **Conv3D Operation:** Describe how a Conv3D kernel differs from a Conv2D kernel in terms of dimensions and movement. Why is this necessary for volumetric data?
- **Memory Challenge:** Conv3D layers are computationally expensive. What engineering trade-offs must be made regarding kernel size, number of filters, or input depth ( $D$ ) when designing a Conv3D architecture?

#### 3.2 Exercise 3: Conv3D Block and Engineering Discipline

##### Instructions:

1. Define a simple Conv3D block in Keras to understand the input and output shapes.
2. Integrate **MLflow tracking** to log the full model configuration and key metrics for the Conv3D block experiment.

```

1 import mlflow
2 import numpy as np
3
4 def simple_conv3d_block(input_shape=(32, 32, 32, 1)):
5     # Simple block for demonstration: D x H x W x C
6     inputs = keras.Input(input_shape)
7
8     # Conv3D layer: 16 filters, 3x3x3 kernel
9     x = keras.layers.Conv3D(16, (3, 3, 3), activation='relu', padding='same')(inputs)
10    x = keras.layers.MaxPool3D((2, 2, 2))(x)
11
12    # TODO: Add a second Conv3D block (32 filters) and another MaxPool3D
13
14    x = keras.layers.Flatten()(x)
15    outputs = keras.layers.Dense(1, activation='sigmoid')(x) # Dummy
16    output
17
18    return keras.Model(inputs, outputs)
19
20 if name == 'main':
21     mlflow.set_experiment("3D_Volumetric_Analysis")
22     with mlflow.start_run(run_name="Conv3D_Baseline"):
23         model_3d = simple_conv3d_block()
24
25         # Log Architecture (Engineering Practice)

```

```

24     model_config = model_3d.to_json()
25     mlflow.log_dict({"model_config": model_config}, "artifacts/
26     model_architecture.json")
27
28     # Log Hyperparameters
29     mlflow.log_param("optimizer", "adam")
30     mlflow.log_param("filters_start", 16)
31
32     # Simulate training and log metrics
33     # TODO: Simulate training by logging a final metric value (e.g.,
34     final_val_loss)
35     # mlflow.log_metric(...)
36     print("MLflow tracking complete for 3D block experiment.")

```

Listing 3: Simple Conv3D Block with MLOps Logging

## 4 Conclusion

### To submit:

- The link to your GitHub repository with the complete and executable `unet_segmentation.py` file, including the Dice and IoU metrics.
- A short report (.pdf or Overleaf link) detailing the results of your segmentation experiment, comparing the convergence speed and final performance (Dice and IoU) of your U-Net model with the concepts introduced in Part 3.